

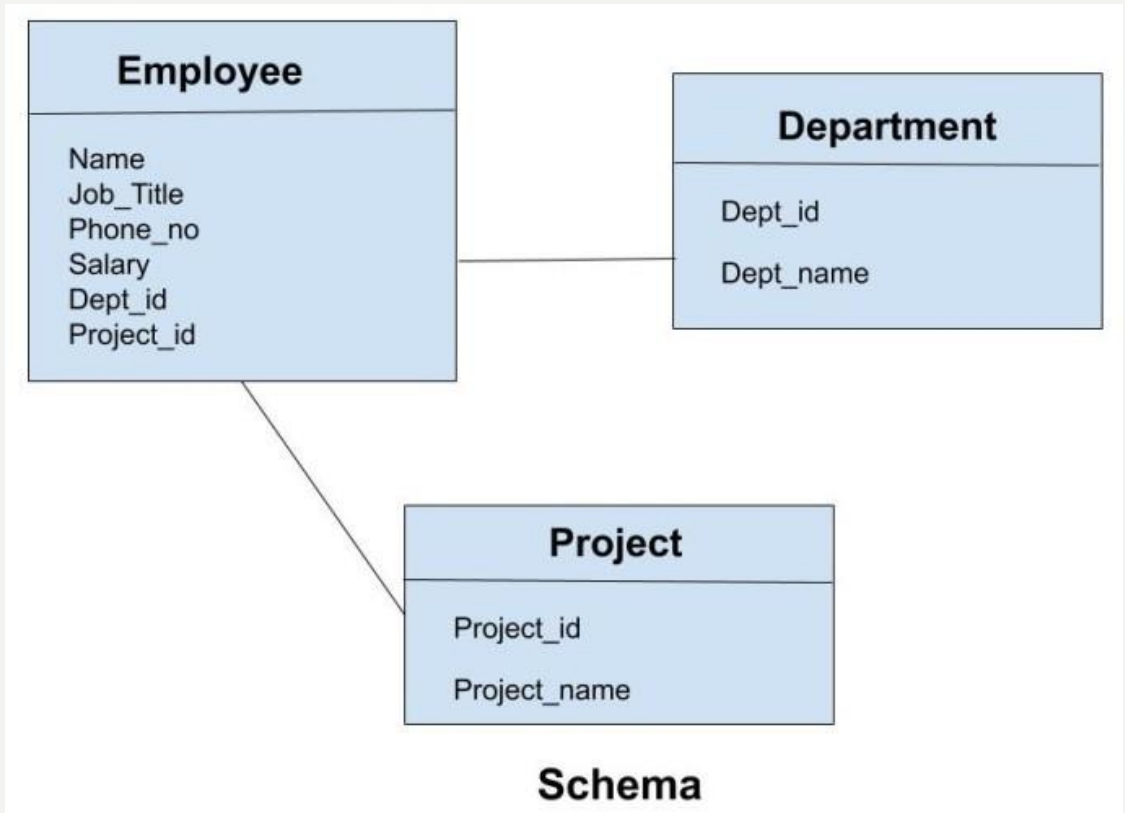


mongoDB[®]

MODULE 5

DATABASE SCHEMA

- A database schema is an abstract design that represents the storage of your data in a database.
- It describes both the organization of data and the relationships between tables in a given database.
- Developers plan a database schema in advance so they know what components are necessary and how they will connect to each other.




NOSQL DATABASE

- (aka "not only SQL") are non-tabular databases and store data differently than relational tables.
- store data in a format other than relational tables.
- come in a variety of types based on their data model.
- The main types are
 - document,
 - key-value,
 - wide-column, and
 - graph.
- They provide flexible schemas and scale easily with large amounts of data and high user loads.

NOSQL DATABASES - NEED


- As storage costs rapidly decreased, the amount of data that applications needed to store and query increased.
- This data came in all shapes and sizes — structured, semi-structured,— and defining the schema in advance became nearly impossible.
- NoSQL databases allow developers to store huge amounts of unstructured data, giving them a lot of flexibility.

SQL DATABASES




Relational


NoSQL DATABASES




Column



Graph



Key-Value



Document

NOSQL VS RELATIONAL DATABASES

Relational databases

- Use tables/rows/columns
- Need a predefined schema / complicated to change
- slow queries when joining multiple tables
- vertically scalable
- Typically closed source

NoSQL

- Originally non-SQL/non-relational
- Not only SQL
- Non-relational databases
- Don't use tables/rows/columns
- Schema-less/easy changes
- Fast queries
- Horizontally scalable
- Open source

DOCUMENT DATABASE - OVERVIEW

Document database

Collection

```
{  
  {doc1}  
  {doc2}  
  .  
  .  
  .  
  {docn}  
}
```

Collection

```
{  
  {doc1}  
  {doc2}  
  .  
  .  
  .  
  {docn}  
}
```

Collection

```
{  
  {doc1}  
  {doc2}  
  .  
  .  
  .  
  {docn}  
}
```

Documents -> rows

Collections -> tables

DOCUMENTS

- Set of key-value pairs
- **Keys:** strings
- **Values:** numbers, strings, booleans, arrays or objects
- **Schemaless:** no need to specify the structure
- **Formats:** JSON, BSON, YAML, or XML

Document Queries

All the users who live in New York and like hiking

All the users older than 40

User's data by user_id

Documents - JSON format

```
{
  "user_id": 512,
  "name": "Carol",
  "last_name": "Harper",
  "email": "carolharper@datazy.com",
  "address": {
    "street": "123 Sesame Street",
    "city": "New York City",
    "state": "New York",
    "country": "USA"
  },
  "hobbies": [
    "hiking",
    "painting"
  ]
}
```


DOCUMENTS – POLYMORPHIC MODEL

```
{
  "user_id": 512,
  "name": "Carol",
  "last_name": "Harper",
  "email": "carolharper@datazy.com",
  "address": {
    "street": "123 Sesame Street",
    "city": "New York City",
    "state": "New York",
    "country": "USA"
  },
  "hobbies": [
    "hiking",
    "painting"
  ]
}
```

```
{
  "user_id": 513,
  "name": "Benjamin",
  "last_name": "Lieberman",
  "email": "benjaminlieberman@datazy.com",
  "date_of_birth": "07/04/1984",
  "hobbies": [
    "reading"
  ]
}
```

COLLECTIONS

- Sets of documents
- Store the same type of entities
- Organize documents and collections by thinking about the queries

POPULAR DOCUMENT DATABASES



DynamoDB



Couchbase



Firebase
Realtime Database

PROS AND CONS OF DOCUMENT DATABASES

Advantages

- Flexibility
 - Don't need to predefine the schema
 - Documents can vary over time
- intuitive for developers
 - Natural way to work
 - JSON is human-readable
 - Less Coding
 - simpler and faster development
 - Easier for new developers
- Horizontal scalability
 - Sharding

Limitations

- more responsibility
 - care about data in the application code (check required email)
 - Care about redundant data (modify duplicated name)

WHEN TO USE DOCUMENT DATABASES

Suitable cases


- **Catalogs**
 - E-commerce web sites/applications store product information
 - Different attributes between the products
 - Embed related information



```
{  
  "product_id": 879,  
  "name": "Fashion shirt",  
  "category": {  
    "category_id": 15,  
    "name": "Tops & t-shirts",  
    "type": "Shirt"  
  }  
}
```

Suitable cases

- **Event logging**
 - Types of events:
 - User logging
 - Product purchase
 - Errors
 - ...
 - Sharding by:
 - Time
 - Type of event
 - ...



```
{  
  "type": "info",  
  "message": "user_logged",  
  "user_id": 551,  
  ...  
}
```

Suitable cases

- User profiles

```
{  
  "user_id": 512,  
  "name": "Carol",  
  "last_name": "Harper",  
  "email": "carolharper@datazy.com",  
  "address": {  
    "street": "123 Sesame Street",  
    "city": "New York City",  
    "state": "New York"  
  },  
  ...  
}
```



- Information may vary
- Document flexibility

Suitable cases

- **Content management systems**
 - Blogs, video platforms, etc.
 - Users' content
 - Comments
 - Images
 - Videos
 - ...



```
{
  "id": 458,
  "url": "myblog/datazy.com",
  "title": "How to write a blog entry at Datazy",
  "tags": [
    "Datazy",
    "Blog"
  ],
  "last10comments": [
    { "name": "Eliza", "comment": "Great!" },
    { "name": "Eric", "comment": "Thank you!" }
  ]...
}
```


Suitable cases

- **Real-time analytics**
 - Page views, unique visitors...
 - Easy to store the information



```
{  
  "_id": "1000241",  
  "hour": "Sat Jun 12 2021 16:40:00 GMT+0200 (EST)",  
  "site": "datazy",  
  "uniques": 5,  
  "pageviews": 15,  
  ...  
}
```

UNSUITABLE CASES

- Very structured data
- Always have consistent data

WHAT IS MONGODB? - OVERVIEW



- MongoDB is a popular document database.
- can be installed locally or hosted in the cloud.
- It stores data in a type of JSON format called BSON.
- A record in MongoDB is a document, which is a data structure composed of key value pairs similar to the structure of JSON objects.

MONGODB - FEATURES

- MongoDB Query Language (MQL)
- `db.users.find({"address.zipcode" : "423201"})`
- Native drivers for programming languages : C#, Java, Python, Scala etc.
- Indexes on any field
- Joins in queries
- scale horizontally
- Replication (50 copies of our data)

MONGODB- PRODUCTS

- **MongoDB Compass:**

- Free GUI
- Explore schema, create queries visually...

- **MongoDB Atlas:**

- Cloud service
- AWS, Azure, Google Cloud

- **MongoDB Enterprise Advanced:**

- Run MongoDB in our infrastructure

- **MongoDB Atlas Lake:**

- Query and analyze data
- AWS S3 and MongoDB Atlas
- MQL

- **MongoDB Charts:**

- Visualizations of the data

- **Realm Mobile Database:**

- Store data locally on iOS or Android

LOCAL VS. CLOUD DATABASE

- MongoDB can be installed locally, which will allow you to host your own MongoDB server on your hardware.
 - You can download and use the MongoDB open source Community Server on your hardware for free.
- MongoDB Atlas is a cloud database platform.
 - This is much easier than hosting your own local database.
 - Sign up for a free MongoDB Atlas account to get started.

INSTALLING MONGODB

- MongoDB Community edition is a free, open-source database that is a popular option for powering modern applications
- Download the installer.
 - Download the MongoDB Community .msi installer from the following link:
https://www.mongodb.com/try/download/community?tck=docs_server
 - In the Version dropdown, select the version of MongoDB to download.
- Run the MongoDB installer.
 - For example, from the Windows Explorer/File Explorer:
 - Go to the directory where you downloaded the MongoDB installer (.msi file).
 - Double-click the .msi file.
 - In the Platform dropdown, select Windows.
 - In the Package dropdown, select msi.
 - Click Download.
- Select **Install MongoDB as a Service**

- MongoDB Atlas
- MongoDB Enterprise Advanced
- MongoDB Community Edition
 - MongoDB Community Server**
 - MongoDB Community Kubernetes Operator
- Tools
- Mobile & Edge

access to advanced functionality such as auto-scaling, serverless instances (in preview), full-text search, and data distribution across regions and clouds. Deploy in minutes on AWS, Google Cloud, and/or Azure, with no downloads necessary.

Give it a try with a free, highly-available 512 MB cluster.

Version

6.0.2 (current)

Platform

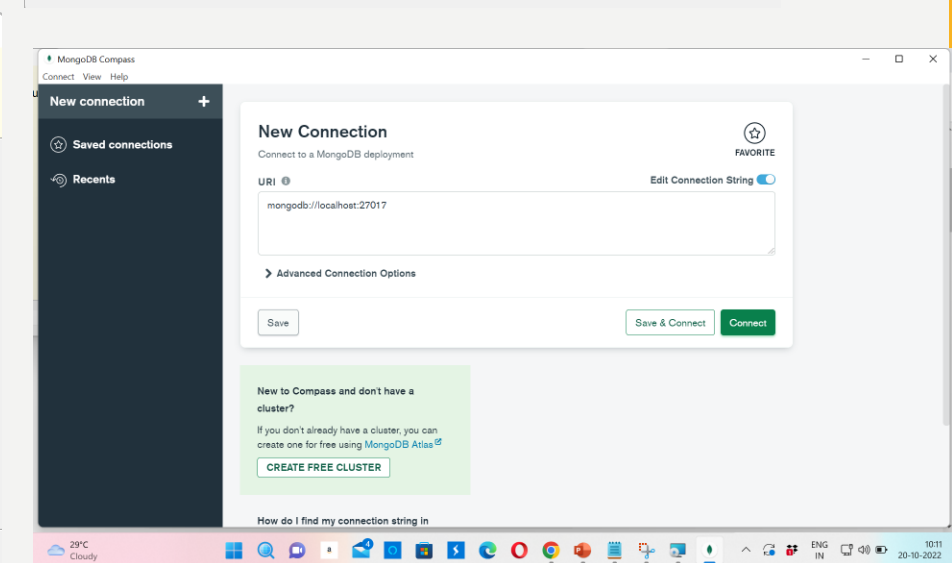

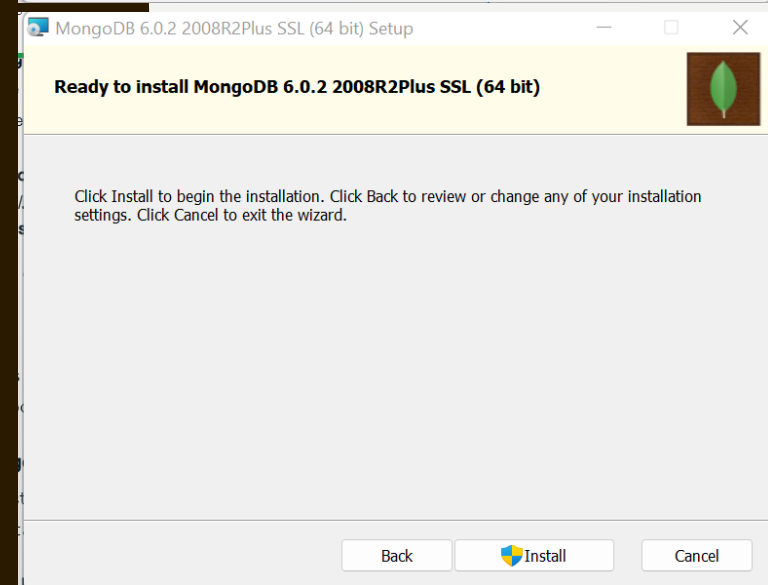
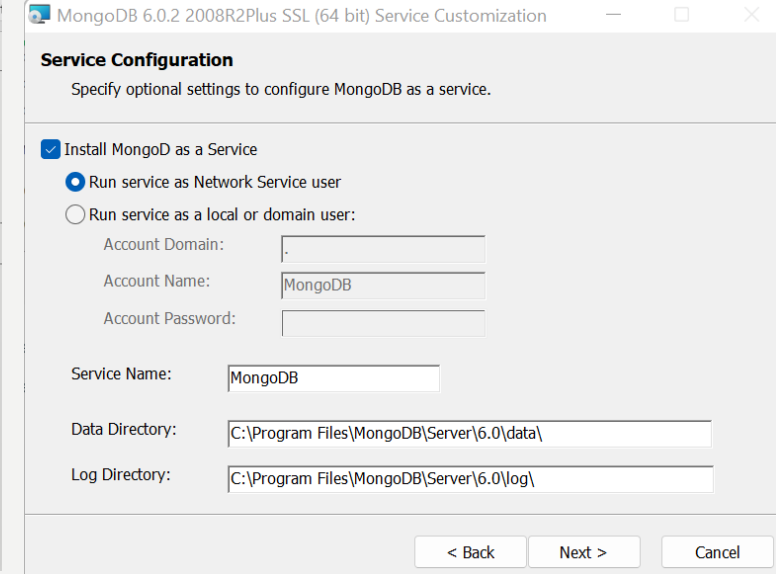
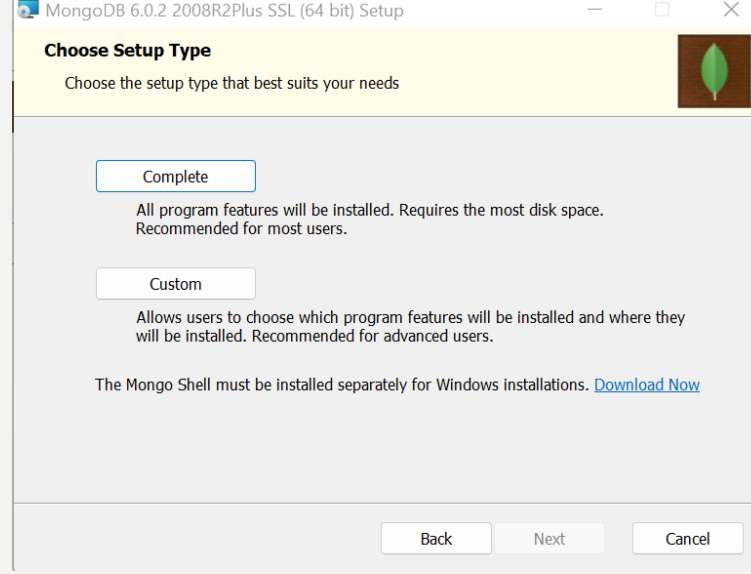
Windows

Package

msi

Download

More Options

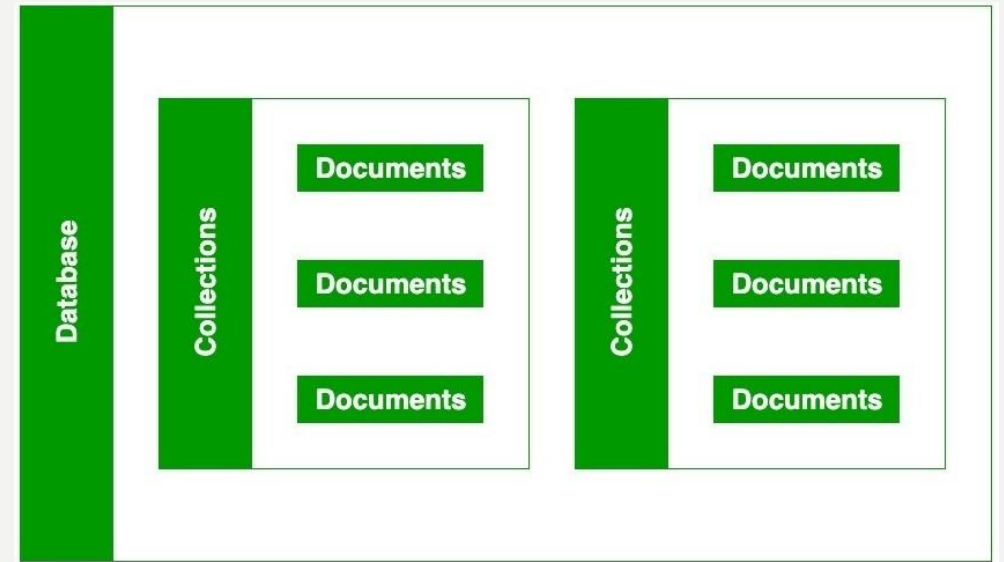


MONGODB COMPASS

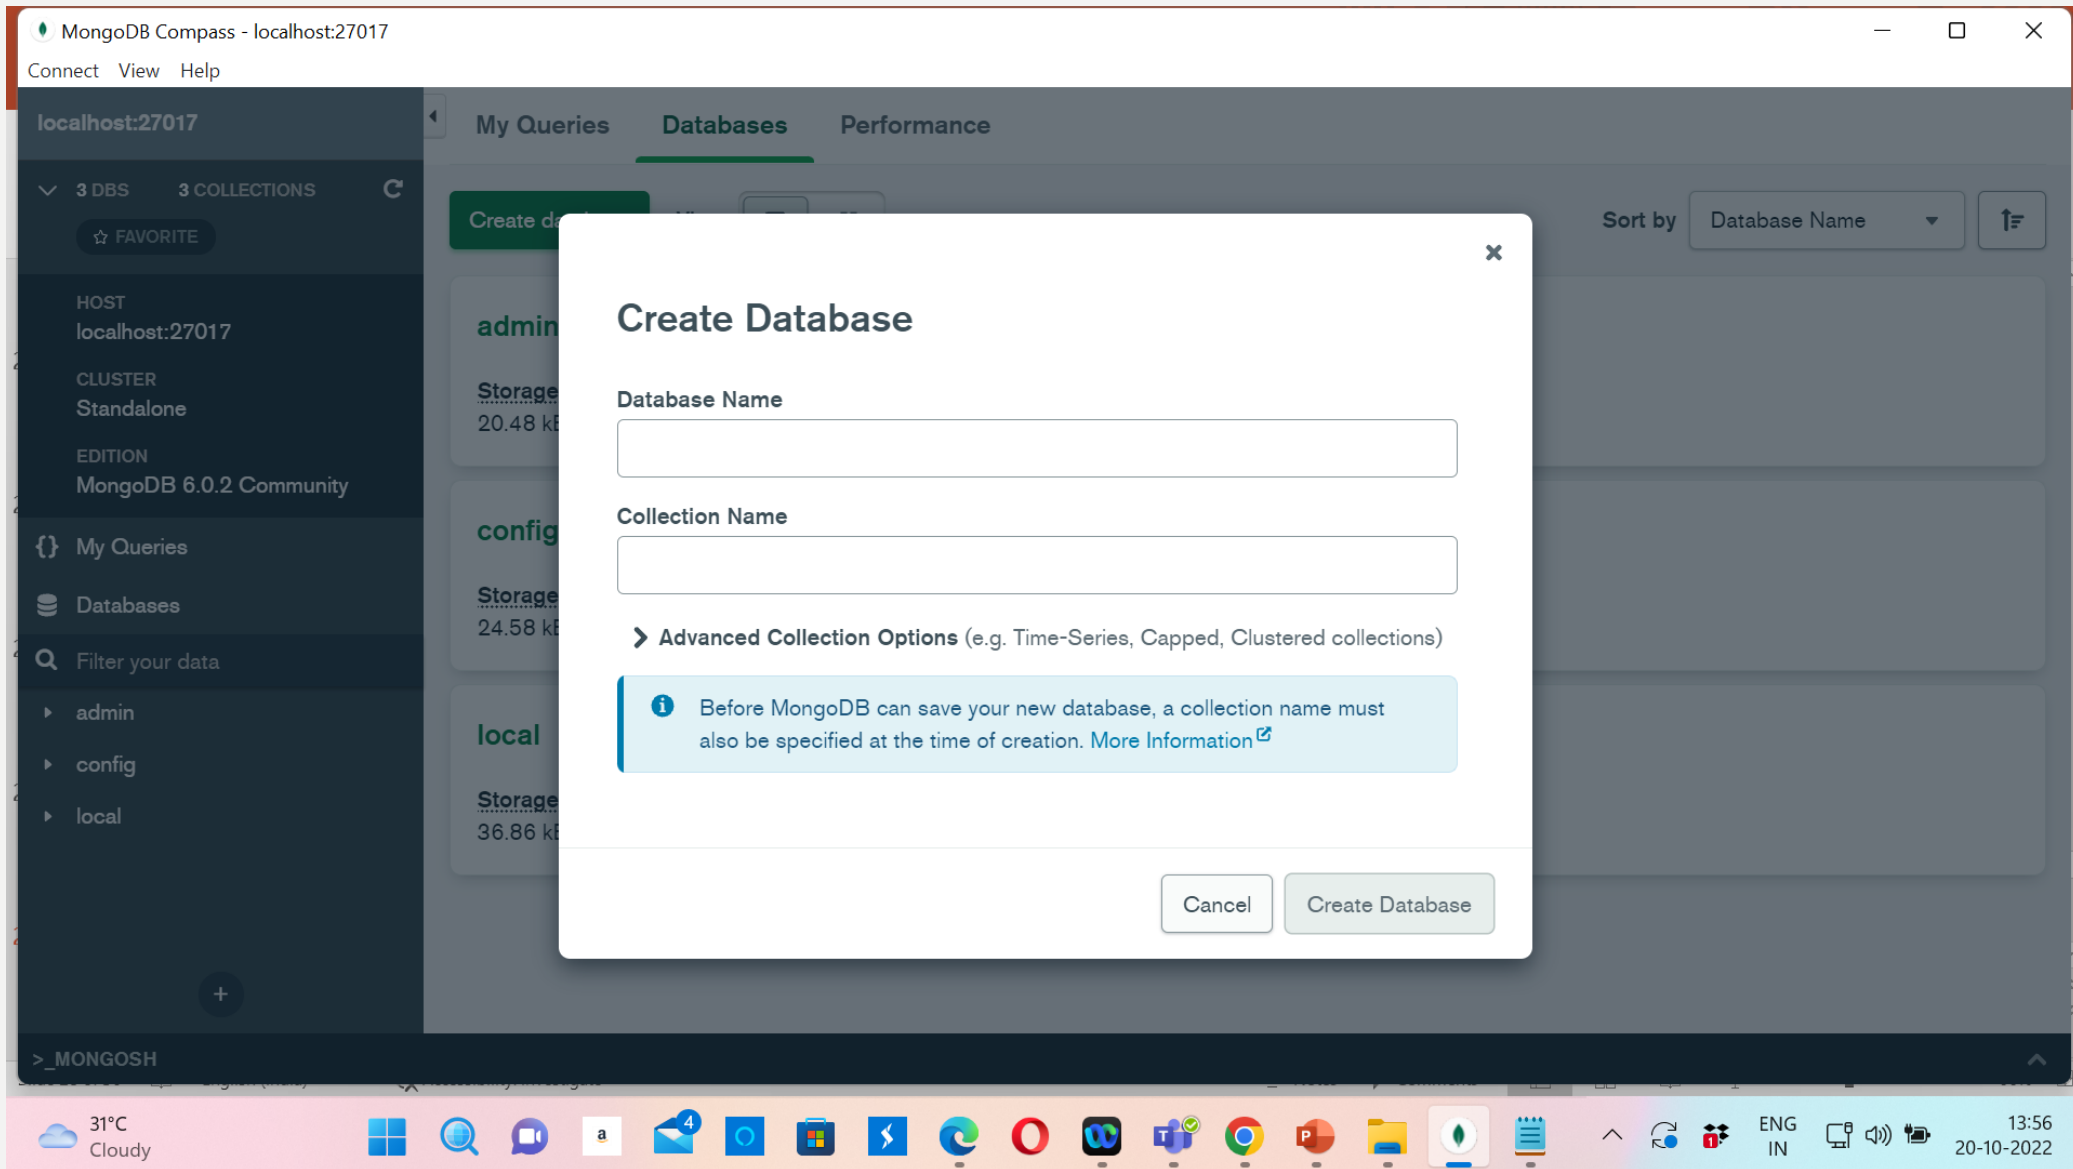
- offers additional functionality like data visualization and performance profiling as well as offering CRUD (create, read, update, delete) access to data, databases, and collections.

HOW MONGODB WORKS?

- MongoDB is structured on a client-server model where a server daemon accepts connections from clients and processes database actions from them. The server must be running for clients to connect and interact with databases.
- Data storage under MongoDB is different from traditional databases. A record in MongoDB is a document (a data structure composed of field and value pairs, similar to JSON objects) and documents are stored in collections (analogous to tables in RDBMS).



CREATING A MONGODB DATABASE



INSERT FIRST DATA

- Create a collection with name books and insert the following data

```
[  
  { "_id" : 8752, "title" : "Divine Comedy", "author" : "Dante", "copies" : 1 },  
  { "_id" : 7000, "title" : "The Odyssey", "author" : "Homer", "copies" : 10 },  
  { "_id" : 7020, "title" : "Iliad", "author" : "Homer", "copies" : 10 },  
  { "_id" : 8645, "title" : "Eclogues", "author" : "Dante", "copies" : 2 },  
  { "_id" : 8751, "title" : "The Banquet", "author" : "Dante", "copies" : 2 }  
]
```

Check out editing, deleting, and find options available on compass

MONGOSH

- MongoDB Shell is the quickest way to connect to (and work with) MongoDB. Easily query data, configure settings, and execute other actions with this modern, extensible command-line interface
- Can be downloaded from
 - <https://www.mongodb.com/try/download/shell2>
- Extract the zip file and update the path



CREATE

READ

UPDATE

DELETE

C R U D

Operations

<https://www.mongodb.com/docs/manual/crud/>

RELATIONSHIP BETWEEN DATABASE, COLLECTION AND A DOCUMENT

Database	humanResourcedb TutorialsTeacher.com		
Collections	employees	addresses	departments
Documents	{ "firstName": "John", "lastName": "King", "email": "john.king@abc.com", "salary": "33000" }	{ "street": "H12", "house": "33", "city": "New York", "country": "USA" }	{ "name": "Technical", "totalEmployees": "100", }

MONGOSH COMMANDS

- **show dbs** – displays all the databases
- **use <database name>** - to change or create to a database
- **db.createCollection(<collection name>)** – to create a collection
- **Insertion**
 - **db.collection_name.insertOne(object)** – to insert an object into a collection
 - **db.collection_name.insertMany()** can be used to insert multiple objects
- **find()** and **findOne()** – to find and select data from MongoDB collection
 - db.collection_name.find()
 - db.collection_name.findone()
 - db.collection_name.find({category:“News”})

```
db.posts.insertOne({  
  title: "Post Title 1",  
  body: "Body of post.",  
  category: "News",  
  likes: 1,  
  tags: ["news", "events"],  
  date: Date()  
})
```

MANGOSH COMMANDS

- Updation
 - `updateOne()` – update the first document that is found matching the provided query.
 - `updateMany()` - update all documents that match the provided query.
- Deletion
 - `deleteOne()` method will delete the first document that matches the query provided.
 - `deleteMany()` method will delete all documents that match the query provided.

IMPORTANT POINTS

- MongoDB reserves `_id` name for use as a unique primary key field that holds ObjectId type.
- A document field name cannot be null but the value can be.
- A document fields can be without quotation marks " " if it does not contain spaces, e.g. { **name**: "Steve"}, { **"first name"**: "Steve"} are valid fields.
- Use the dot notation to access array elements or embedded documents.
- MongoDB supports maximum document size of 16mb.
- MongoDB collection can store documents with different fields. It does not enforce any schema.

UPDATE OPERATORS

- The following operators can be used to update fields:
 - **\$inc**: Increments the field value
 - **\$rename**: Renames the field
 - **\$set**: Sets the value of a field
 - **\$unset**: Removes the field from the document

- `updateMany()` : It update all documents in a collection with matching filter.
- `updateOne()` : It update only one top most document in a collection with matching filter.
- `update()` : By default, the `update()` method updates a single document. Include the option `{multi : true}` to update all documents that match the query criteria. Hence we can use it as both ways.

```
Syntax : db.collection.update(  
  <query>,  
  <update>,  
  {  
    upsert: <boolean>,  
    multi: <boolean>,  
    writeConcern: <document>,  
    collation: <document>  
  }  
)
```

MONGODB QUERY OPERATORS

Comparison Operators

\$eq : Values are equal
\$ne : Values are not equal
\$gt : Value is greater than another value
\$gte : Value is greater than or equal to another value
\$lt : Value is less than another value
\$lte : Value is less than or equal to another value
\$in : Value is matched within an array

Logical Operators

\$and : Returns documents where both queries match
\$or : Returns documents where either query matches
\$nor : Returns documents where both queries fail to match
\$not : Returns documents where the query does not match

PROJECTION

Project Fields to return from Query

- queries in MongoDB return all fields in matching documents by default.
- A projection document is used to restrict fields to return.
- `db.collection.find()` returns all fields in matching documents if you do not specify a projection document.
 - `db.inventory.find()`
 - `db.inventory.find({ status: "A" })`
- Return specific fields and the `_id` field only
 - A projection can explicitly include several fields by setting the `<field>` to `1` in the projection document.
 - `db.inventory.find({ status: "A" }, { item: 1, status: 1 })`
- suppress `_id` Field
 - remove the `_id` field from the results by setting it to `0` in the projection
 - `db.inventory.find({ status: "A" }, { item: 1, status: 1, _id: 0 })`
- Return All But the Excluded Fields
 - `db.inventory.find({ status: "A" }, { status: 0, instock: 0 })`

EMBEDDED DOCUMENTS

- A document in MongoDB can have fields that hold another document. It is also called nested or embedded document.
- Example is shown right
 - Department and address field contains another document
 - the address field contains the phone field which holds a second level document.

Example: Embedded Document

```
{
  _id: ObjectId("32521df3f4948bd2f54218"),
  firstName: "John",
  lastName: "King",
  department: {
    _id: ObjectId("55214df3f4948bd2f8753"),
    name: "Finance"
  },
  address: {
    phone: { type: "Home", number: "111-000-000" }
  }
}
```


EMBEDDED DOCUMENTS

- An embedded document can contain upto 100 levels of nesting.
- Supports a maximum size of 16 mb.
- Embedded documents can be accessed using dot notation `embedded-document.fieldname`, e.g. access phone number using `address.phone.number`.

QUERYING THE EMBEDDED DOCUMENTS

- Matching an embedded or nested document
 - `db.inventory.find({ size: { h: 14, w: 21, uom: "cm" } })`
- query on nested field
 - `db.inventory.find({ "size.uom": "in" })`
- specify the match using query operator
 - `db.inventory.find({ "size.h": { $lt: 15 } })`

ARRAY

- Arrays can hold any type of data or embedded documents.
- A field in a document can hold array.
- Array elements in a document can be accessed using dot notation with the zero-based index position.
- The example document contains the skills field that holds an array of strings. To specify or access the second element in the skills array, use skills.1.

Example: MongoDB Document with an Array

```
{
  _id: ObjectId("32521df3f4948bd2f54218"),
  firstName: "John",
  lastName: "King",
  email: "john.king@abc.com",
  skills: [ "Angular", "React", "MongoDB" ],
}
```

QUERYING THE ARRAY

- Create a collection with the name :inventory2
- Insert the following data
 - db.inventory2.insertMany([
 - { item: "journal", qty: 25, tags: ["blank", "red"], dim_cm: [14, 21] },
 - { item: "notebook", qty: 50, tags: ["red", "blank"], dim_cm: [14, 21] },
 - { item: "paper", qty: 100, tags: ["red", "blank", "plain"], dim_cm: [14, 21] },
 - { item: "planner", qty: 75, tags: ["blank", "red"], dim_cm: [22.85, 30] },
 - { item: "postcard", qty: 45, tags: ["blue"], dim_cm: [10, 15.25] }
 -]);

QUERYING THE ARRAY

- Match an Array

- queries for all documents where the field tags value is an array with exactly two elements, "red" and "blank", in the specified order:
- `db.inventory2.find({ tags: ["red", "blank"] })`
- to find an array that contains both the elements "red" and "blank", without regard to order or other elements in the array, use the \$all operator:
- `db.inventory2.find({ tags: { $all: ["red", "blank"] } })`

- Query an Array for an Element

- To query if the array field contains at least one element with the specified value, use the filter { <field>: <value> } where <value> is the element value.
- `db.inventory2.find({ tags: "red" })`

- To specify conditions on the elements in the array field, use query operators in the query filter document:

- { <array field>: { <operator I>: <value I>, ... } }
- `db.inventory2.find({ dim_cm: { $gt: 25 } })`

DATATYPES ALLOWED IN SCHEMAS / SCHEMA TYPES

- String
- Number
- Date
- Boolean
- Buffer
- ObjectId
- Mixed
- Array

FETCHING DATA FROM STRUCTURED FILES

- Data can also be updated from csv file

Import To Collection dw28.PlayTennis

Select File

Select a file...

Select Input File Type

JSON

CSV

Options

☐ Stop on errors

CANCEL

IMPORT

RELATIONSHIP BETWEEN DATA

- In MongoDB,
 - one-to-one,
 - one-to-many, and
 - many-to-many
- relations can be implemented in two ways:
 1. Using embedded documents
 2. Using the reference of documents of another collection

IMPLEMENTING RELATION

- Related data can be included using Embedded Document
- For example, you can include an address as an embedded document, as shown below.

```
db.employee.insertOne({
  _id: ObjectId("32521df3f4948bd2f54218"),
  firstName: "John",
  lastName: "King",
  email: "john.king@abc.com",
  salary: "33000",
  DoB: new Date('Mar 24, 2011'),
  address: {
    street: "Upper Street",
    house: "No 1",
    city: "New York",
    country: "USA"
  }
})
```

- Implement Relation using Reference is another way
- using the reference of the primary key field of documents of another collection.

```
db.address.insertOne({
  _id: 101,
  street: "Upper Street",
  house: "No 1",
  city: "New York",
  country: "USA"
})

db.employee.insertOne({
  firstName: "John",
  lastName: "King",
  email: "john.king@abc.com",
  salary: "33000",
  DoB: new Date('Mar 24, 2011'),
  address: 101
})
```

ONE TO ONE USING EMBED METHOD

```
{  
  _id: "joe",  
  name: "Joe Bookreader",  
  address: {  
    street: "123 Fake Street",  
    city: "Faketon",  
    state: "MA",  
    zip: "12345"  
  }  
}
```

ONE TO ONE USING REFERENCE METHOD

```
// patron document
{
  _id: "joe",
  name: "Joe Bookreader"
}

// address document
{
  patron_id: "joe", // reference to patron document
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}
```

ONE TO MANY

```
// patron document
{
  _id: "joe",
  name: "Joe Bookreader"
}

// address documents
{
  patron_id: "joe", // reference to patron document
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}

{
  patron_id: "joe",
  street: "1 Some Other Street",
  city: "Boston",
  state: "MA",
  zip: "12345"
}
```

```
{
  "_id": "joe",
  "name": "Joe Bookreader",
  "addresses": [
    {
      "street": "123 Fake Street",
      "city": "Faketon",
      "state": "MA",
      "zip": "12345"
    },
    {
      "street": "1 Some Other Street",
      "city": "Boston",
      "state": "MA",
      "zip": "12345"
    }
  ]
}
```

CONNECTING TO MONGODB FROM PYTHON

- PyMongo is a Python distribution containing tools for working with MongoDB, and is the recommended way to work with MongoDB from Python.
- To work with pymongo, first install pymongo distribution
 - `pip install pymongo`
- To check whether the library is properly installed or not, run
 - `import pymongo`
- Making a connection with a MongoClient
 - create a MongoClient instance to run the mongodb instance
 - `from pymongo import MongoClient`
 - `client = MongoClient()`
(or)
 - `client = MongoClient('localhost', 27017)`

PYMONGO

- Getting the database names
 - `client.list_database_names()`
- Getting a database
 - A single instance of MongoDB can support multiple independent databases.
 - `db = client.test_database`
(or)
 - `db = client['test-database']`
- Getting the collection names
 - `db.list_collection_names()`
- inserting documents into collection
 - `insert_one()` - used to insert a document.
 - `insert_many()` - used to insert many documents.
- Querying
 - `find_one()` - is used to get a single document
 - `find()` - returns all the documents
- `count_documents()` - to find how many documents match a query

```
>>> db.test.insert_one({'x': 1})
<pymongo.results.InsertOneResult object at 0x0000018A9EA06EE0>
>>> db.test.insert_one({'y': 2})
<pymongo.results.InsertOneResult object at 0x0000018A9D9F4B20>
>>> db.test.find_one()
{'_id': ObjectId('635e877308d42c347b2c4256'), 'x': 1}
>>> db.test.find()
<pymongo.cursor.Cursor object at 0x0000018A9E9A0CA0>
>>> for x in db.test.find():
...     print(x)
...
{'_id': ObjectId('635e877308d42c347b2c4256'), 'x': 1}
{'_id': ObjectId('635e878408d42c347b2c4257'), 'y': 2}
```

```
>>> db.inventory.count_documents({})
5
>>> db.inventory.count_documents({"status": "A"})
3
```

CHECK YOUR UNDERSTANDING

- MongoDB is a (relational / NoSQL) database.
- _____ is the command-line client of MongoDB Database? (Mongo Shell / Web application/ windows service)
- which command starts mongo shell? (mongo / mongod / mongosh / shell)
- Which command creates a new database? (create / switch / use / new)
- MongoDB collections are similar to ____ in RDBMS. (Views / Tables / Rows / Columns)
- Which method creates a new collection?
- MongoDB stores data as _____ document (JSON/ BSON/ XML/YML)